

# COMP 474/6741 Intelligent Systems (Winter 2024)

## Worksheet #9: Introduction to Deep Learning

**Task 1.** Let's compute the *cross-entropy loss* for a binary classification task, where a neural network has to classify an image into *hot dog* (label 1) vs. *not hot-dog* (label 0), using the loss function

$$H(y, p) = -(y \ln(p) + (1 - y) \ln(1 - p))$$

Assume for a given input image, we obtained an output value  $p$  (using the sigmoid activation function) of 0.84. The expected output  $y$  for this image is 1. Compute the cross-entropy loss in this case: .....

What's the loss if the network had predicted 0.2 instead (same image, expected result is still 1): .....

**Task 2.** Ok, let's now compute the *categorical cross-entropy loss* for a classification task where a network has to distinguish between three classes: *cat*, *dog*, and *hot-dog*. The network outputs a probability distribution across these classes for a given image, using the *softmax* activation function in the output layer. The loss function is defined as:

$$L(y, p) = - \sum_{i=1}^N y_i \ln(p_i)$$

For a given image, the network predicted the following probabilities: *cat* = 0.2, *dog* = 0.5, *hot-dog* = 0.3. The true class of the image is *dog*, represented as a one-hot encoded vector  $y = [0, 1, 0]$ .

Calculate the categorical cross-entropy loss for this scenario: .....

**Task 3.** Consider the following matrix that represents an image. This image will be fed into a convolutional neural network (CNN):


|   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|
| 1 | 1 | 2 | 2 | 2 | 0 | 0 |
| 2 | 0 | 1 | 1 | 2 | 1 | 2 |
| 0 | 1 | 0 | 0 | 1 | 1 | 2 |
| 0 | 2 | 1 | 2 | 0 | 2 | 2 |
| 1 | 2 | 0 | 0 | 1 | 0 | 1 |
| 0 | 0 | 0 | 0 | 1 | 2 | 1 |
| 2 | 0 | 0 | 0 | 2 | 1 | 1 |

Assume that we use the following convolution filter with a stride of 2 (no padding):

|   |    |    |
|---|----|----|
| 0 | 1  | 1  |
| 0 | 1  | 0  |
| 0 | -1 | -1 |

What will be the size of the activation map? .....

What will be the resulting activation map?



**Task 4.** The task here is to understand the structure of the CNN (shown at the top) we are building: Our first convolution layer has 10 filters, each sized  $4 \times 4 \times 3$  (`kernel_size = (4,4)`), thus  $(4 \times 4 \times 3 + 1) \times 10 = 490$  weights (parameters to train).

How do we obtain the output shape of this layer? The general formula you can use is (for `padding = "same"`, meaning the size of the kernel is the same as the input, padded with zeros):<sup>1</sup>

$$\text{output\_shape} = \left( \text{None}, \frac{\text{input\_height}}{\text{stride}}, \frac{\text{input\_width}}{\text{stride}}, \text{filters} \right)$$

Now compute the output tensor shape of the first convolution layer using the formula above:

output\_shape = .....

In the second convolution layer, we want to apply 20 filters of size  $3 \times 3$  and a depth of 10.

How many weights do we have to train? .....

Compute again the output tensor shape using the formula above:

output\_shape = .....

**Task 5.** What will be the output of a pooling layer with a size of  $2 \times 2$  and a stride of 1, on the activation map of Task 3 above, if we use the following strategies:

1. Average pooling:

2. Max pooling:

<sup>1</sup>In a Keras model input shape, "None" represents a flexible batch size dimension, allowing for input batches of any size.